

NVLink all-to-all is non-monotonic beyond 16 GPU — placement noise, not an N-scaling law

