

E2E scaling: GB300 stays flat over NVLink; H100 drops to 75% at 16 GPU (2 IB nodes)

