


# GB300 v26.06 — the other two axes: GPUs cost ~1-3% per 4x (cause unresolved); cache ratio buys no throughput

## Weak scaling: the token-work is flat; c is a fit artifact



## Cache ratio is a throughput no-op: 3.3-3.8x HBM, no trend

